

# Hadoop

## Do you need a relational SQL DB?

1. difficult for large scale data ("planet-wide")
2. fast when de-normalized
3. transactions
4. export from hadoop to mysql?
5. Example Infrastructure
6. Data Source -> Spark Streaming -> MongoDB -> Web FrontEnd
7. "CAP Theorum"
8. Consistency
  1. diff. profile thumbs per-request
9. Availability
10. Partition Tolerance <- non-negotiable for big data

## History

- google releases paper -> open source implementation
- typically apache projects made with java
- two major pieces to hadoop
  1. hdfs – stores files in folders
  2. 64MB to 2GB standardized chunks
  3. three replicas of each block
  4. blocks distributed across many machines
  5. very poor storage efficiency!
    - bad *storage*...
  6. mapreduce
    - first execution paradigm
    - paper released a year later

## MapReduce

```
* map, reduce -- two stages
* share nothing
* insensitive to order
```

## Mappers

```
* filter, transform, parse
* projection to (key, value)
```

## Reducers

- groups by mapper's key
- aggregating step
- counts, statistics
- joins

## Hadoop Ecosystem

- "Using Hadoop" cf. "Using Windows"
- Hadoop provides an operating system for data
  - storage + processing
- Higher-level tools are layered on top
- HBase
- Accumulo
- Storm
- Spark
- Zookeeper, Kafka, ...
- Each hadoop installation / project is different

## Advantages of Hadoop

- Linear Scaling
  - 2x computers => c. 2x speed
  - 2x data => 2x slower
  - general linear heuristics *work*
- Schema on Read
  - ETL, Relational, Schema'd Systems
    - Schema on-write
    - Design up-front/
  - Hadoop
    - Schema on-read
    - Keep *original* data around
    - Design at solution-time, ie., late
      - multiple views/structures, depending on *query*
    - parse on-read
    - unstructured data (logs, audio, images...)
- Automatic (transparent) Parallelism
  - mrjob, pig job, etc. are automatically parallel
  - do not need to consider:
    - threading, locking, fault tolerance, etc.
  - as long as you follow the hadoop paradigm
    - ie., map-reduce
  - strikes the right balance of abstraction
    - ...or not?

# Limitations

- \* Cannot edit files
  - \* "replace", ie., delete & copy
- \* Append
- \* Applications create files
  - \* files are written to hdfs
- \* Getting data into hadoop is its own engineering problem
  - \* data flow systems ususally needed

## Python and Hadoop

- Python vs. Java
  - Why python?
    - data science
    - scripting
    - ease of use
    - ecosystem
    - community
  - Why not java?
    - compiling is its own engineering effort
    - dependency (jar) hell
    - sensitivity to build enviornment
  - Why not python?
    - with hadoop, slower / needs translation

## MrJob

- Hadoop streaming
  - abitary programs that use stdin/out
  - stdin/out is slower than jvm in-memory data-passing (ie. API methods)
    - multi-jvm
      - jvm/core (/pc)
- Integrates with Amazon EMR or Hadoop
  - Elastic Map-Reduce

Using open source tools such as Apache Spark, Apache Hive, Apache HBase, Apache Flink, and Presto, coupled with the dynamic scalability of Amazon EC2 and scalable storage of Amazon S3, EMR gives analytical teams the engines and elasticity to run Petabyte-scale analysis for a fraction of the cost of traditional on-premise clusters.

## The Map Reduce Paradigm

- Aside: generators and `yield`

- Generators are functions with “suspended returns”
  - They calculate values on-demand, but behave like lists
    - ie., they are iterable
- They are essentially rewritten into objects with a `next()` method
  - and *internal state* that provides a “current value”
  - rather than `return`, `yield` is read as “when next() is called”
  - the internal state keeps track of looping
  - `.__next__()` calculates the on-demand value
- `next()` will move the generator forward
- `tuple()`, `list()`, ... will *force* the generator
  - causing it to calculate all values and store as a list ```python

```
def numbers(): for i in range(3): yield 3
```

```
n = numbers() ns = [numbers(), numbers()]
```

```
print(next(n)) print(next(n)) print([ tuple(e) for e in ns]) ```
```

## A Map Reduce Example

```
from functools import reduce

def mapper(line):
    for word in line.split(" "):
        yield (word.lower(), 1)

def reducer(pair):
    key, value = pair
    yield (key, sum(value))

def _combine(data, values):
    for (k, v) in values:
        data[k].append(v)
    return data
```

- concerns
  - memory
    - how big is `line` ?
  - mrjob runs in bash
    - system memory management
    - system error managment
    - c. 2x+ slower than java
      - but consider compile time

```

data = ["hello world", "goodbye world"]

# map-side
mapped = map(mapper, data)
combined = reduce(_combine, mapped, defaultdict(list))

# reduce-side
reduced = map(reducer, combined.items())

# report
print([next(result) for result in reduced])

```

::: notes ```python

# or, just:

```

print([]) ```
::

```

## HBase

1. NoSQL database on top of HDFS
2. Based on BigTable (Open Source Version)
3. cf. Google's published papers
4. CRUD operations
5. Region Servers
6. (aprox. Shards = partitions of distributed data)
7. On top of HDFS (itself distributed)
8. HMaster is master over all regions
  1. Zookeeper Orchestrates
9. Data Model
10. Rows referenced by unique KEY
11. Rows have "COLUMN FAMILIES"
  1. with arbitrary number of columns
  2. helpful for sparse data
12. Rows have "CELL"s – row/col intersection
  1. versioned by timestamp
13. Example: Google's Problem
  1. com.cnn.www <- lexographic key choice min read
  2. "Contents" column family with single column
    1. so that content cells are versioned
  3. "Anchor" column family
    1. Contains may "Anchor" columns – eg., millions

- 2. One per url linking
- 14. HBase APIs
  - 1. REST
    - 1. Open Port on VM
    - 2. "starbase" python client
  - 2. Spark, Hive, Pig, ...
- 15. Example: User Ratings
  - 1. RowID = User ID
  - 2. Column Family = {Rating:, Rating:, Rating:,...}
- 16. starbase with python
  - 1. Limited: python is in-memory, not bigdata
  - 2. .create() creates column family
  - 3. .drop(), .close(), ...
  - 4. .batch() (~connection)
    - 1. .update() (~insert)
- 17. Interactive shell
  - 1. create 'users'
  - 2. list
  - 3. scan 'users'
  - 4. disable 'users'; drop 'users'
- 18. Pig
  - 1. Big Data
  - 2. create table upfront, unique keys, etc.
  - 3. hbase://
  - 4. USING ..HBaseStorage

## Cassandra

- 1. Like HBase but no master node
- 2. HBase has HMaster, Zookeeper
- 3. CQL – Cassandra Query Language
- 4. Non-relational
- 5. No Joins
- 6. All queries on primary key (or secondary)
- 7. Shell (CQLSH)
- 8. Eventually Consistent
- 9. CAP – compromise on C
- 10. Fast Access to Rows
- 11. Two Systems (OLTPish and OLAPish)
- 12. Ring Architecture
- 13. ( Region ID ranges in Ring )
- 14. Gossip Protocol (negotiation between nodes)
- 15. Nodes share data, gossip to find out which has it
- 16. DataStax = Spark + Cassandra

17. Cassandra appears as a DataFrame

## Pig

1. PigLatin Scripts
2. PigView in admin console
3. Grunt prompt
4. Data Analysis without writing Mappers & Reducers
5. Can be faster than MapReduce with TEX
6. Spark is preferred, Pig not due to performance
7. Historical: Fixed with TEX
8. On top of MapReduce & TEZ
9. TEZ = Directed Acyclical Graph for Optimizing jobs
10. Pig can go via TEZ or MR
11. Example:
12. "relation" = variable = data set
13. "as" provides schema
  1. schemaless default expectation
14. expects tab delimited by default
15. FOREACH relation GENERATE new-schema
16. transformation
17. GROUP relation BY field
18. "bags" data
19. DESCRIBE
20. dumps relation structure
21. FILTER relation BY test
22. JOIN relation BY fields BY relation
23. joins relations
24. renames fields, including full path from original
25. ORDER BY
26. IMPORT, DEFINE, REGISTER
  1. interfacing with user-defined functions (JVM)

## Spark

1. In-fashion
2. "a fast and general engine for data processing"
3. vs. Pig
4. Rich Ecosystem (eg., Machine Learning)
5. Same Pattern as Hadoop
6. Driver PRogram -> Manager -> Cache
7. Can use hadoop, doesn't need to
8. In-Memory Processing System
9. vs. MapReduce - File System

10. ~(10 to 100)x F.Sys
11. Built in TEZ-like optimization system
12. One Concept: "Resilient Distributed Dataset (RDD)"
13. "Data set" add in later version, more SQL-like
14. Libraries Included
15. Spark Streaming
  1. Realtime Analysis (vs., Batch)
16. Spark SQL
  1. SQL interface to Spark
  2. Heavy optimization work (>2.0)
17. MLlib
  1. Machine Learning Lib
  2. vs., eg., with MapReduce (hard to do ML)
18. GraphX
  1. (Social) graph analytics
19. Written in Scala
20. Python Libs available
21. Compiles to bytecode – always faster than python
22. Programming model fits spark more naturally
  1. eg., data transformation via ann. fns.
23. RDDs
  1. SparkContext creates RDDs
  2. Creating
  3. .parrellize([data])
  4. .textFile
    1. hdfs://
  5. HiveContext, Cassandra, ElasticSearch, ..
  6. Transforming
  7. FilterMonadic: .map, .flatMap, .filter
  8. take, top, reduce, count, etc.
  9. Lazy Evaluation
  10. graph of dependent actions built up
  11. nothing happens until "action" is called
  12. Example: Find Lowest Rating
  13. Spark 1 – RDD Interface
    1. MapReduce-like
    2. .map, .reduceByKey, .mapValues, .sortBy, .take
  14. Ambari / Manager
    1. Spark Log Level (INFO -> ERROR)
  15. Spark SQL
  16. Extends RDD to DataFrame
    1. cf. R, Pandas, ...
  17. Data Frames



1. from Spark 2.0, lingua franca across libs
2. DataSet of Row Objects
3. SQL Queries
4. Schema
5. Read/Write to Json, etc.
18. Create DataFrame
  1. from json
  2. from HiveContext
19. DataFrame
  1. .sql
  2. .select, .filter, .mean, .groupBy, ...
  3. .rdd() extracts to RDD level
  4. user-defined functions
20. Example
  1. SparkSession
    1. provides SparkContext
    2. .getOrCreate() recovery
21. MLLib with Spark
22. Recommendation Example
  1. ALS alg. – netflix-prize recommendation alg.
    1. Predict rating for given user with ratings history
  2. from pyspark.ml.recommendation import ALS
  3. read from hdfs://
    1. .cache() in memory
  4. plug in data into model.transform & predict

## Hive

1. SQL on top of MapReduce & TEZ
2. builds on existing SQL knowledge
3. interactive prompt
4. Easy OLAP
5. ie., long-time processing queries
6. Not good for real-time analytics (ie., OLTP)
7. “few minutes” but massive data sets
8. Pretend relational, HDFS is schemaless
9. No inserts/deletes, etc.
10. Pig/Spark more powerful
11. MySQL-like
12. VIEWS
13. Example
  1. HiveView
  2. DROP TABLE
  3. Upload Table, Tab file + Pipe File

4. CREATE VIEW / SELECT / GROUP BY / JOIN, etc

14. "Schema on Read"

1. unstructured data -> structured as read
2. metastore holds schema
3. CREATE TABLE
4. ROW FORMAT
5. FIELDS TERMINATED BY
6. STORED AS
7. OVERWRITE INTO TABLE
8. LOAD DATA, LOAD DATA LOCAL (copy)
9. No relational DB, just parsing structure
10. Managed Tables – Hive Owned, vs., External
11. DROP'able, etc.
12. PARTITIONS
13. sub dirs
14. significant optimization when relevant